

# Advanced SQL



Lecture #02



Database Systems  
15-445/15-645  
Fall 2018

AP

Andy Pavlo  
Computer Science  
Carnegie Mellon Univ.

# RELATIONAL LANGUAGES

---

User only needs to specify the answer that they want, not how to compute it.

The DBMS is responsible for efficient evaluation of the query.

→ Query optimizer: re-orders operations and generates query plan

# SQL HISTORY

---

Originally “**SEQUEL**” from IBM’s **System R** prototype.

- Structured English Query Language
- Adopted by Oracle in the 1970s.

**Original name SEQUEL**

IBM releases DB2 in 1983.

ANSI Standard in 1986. ISO in 1987

- Structured Query Language



# SQL HISTORY

---

Current standard is **SQL:2016**

- **SQL:2016** → JSON, Polymorphic tables
- **SQL:2011** → Temporal DBs, Pipelined DML
- **SQL:2008** → TRUNCATE, Fancy ORDER
- **SQL:2003** → XML, windows, sequences, auto-generated IDs.
- **SQL:1999** → Regex, triggers, OO

**Even though there is a standard, nobody actually follows it**

Most DBMSs at least support **SQL-92**

- System Comparison: <http://troels.arvin.dk/db/rdbms/>

# RELATIONAL LANGUAGES

---

Data Manipulation Language (DML)

Data Definition Language (DDL)

Data Control Language (DCL)

**security authorization**

Also includes:

- View definition
- Integrity & Referential Constraints
- Transactions

Important: SQL is based on **bags** (duplicates) not **sets** (no duplicates).

# TODAY'S AGENDA

---

Aggregations + Group By  
String / Date / Time Operations  
Output Control + Redirection  
Nested Queries  
Common Table Expressions  
Window Functions



# EXAMPLE DATABASE

**student(sid, name, login, gpa)**

sid	name	login	age	gpa
53666	Kanye	kayne@cs	39	4.0
53688	Bieber	jbieber@cs	22	3.9
53655	Tupac	shakur@cs	26	3.5

**course(cid, name)**

cid	name
15-445	Database Systems
15-721	Advanced Database Systems
15-826	Data Mining
15-823	Advanced Topics in Databases

**enrolled(sid, cid, grade)**

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53655	15-445	B
53666	15-721	C

# AGGREGATES

---

Functions that return a **single value** from a bag of tuples:

- **AVG(col)** → Return the average col value.
- **MIN(col)** → Return minimum col value.
- **MAX(col)** → Return maximum col value.
- **SUM(col)** → Return sum of values in col.
- **COUNT(col)** → Return # of values for col.





# AGGREGATES

Aggregate functions can only be used in the **SELECT** output list.

*Get # of students with a “@cs” login:*

```
SELECT COUNT(login) AS cnt  
FROM student WHERE login LIKE '%@cs'
```

```
SELECT COUNT(*) AS cnt  
FROM student WHERE login LIKE '%@cs'
```

```
SELECT COUNT(1) AS cnt  
FROM student WHERE login LIKE '%@cs'
```

# MULTIPLE AGGREGATES

*Get the number of students and their average GPA that have a “@cs” login.*

```
SELECT AVG(gpa), COUNT(sid)
FROM student WHERE login LIKE '@cs'
```

AVG(gpa)	COUNT(sid)
3.25	12

# DISTINCT AGGREGATES

**COUNT**, **SUM**, **AVG** support **DISTINCT**

*Get the number of unique students that have an “@cs” login.*

```
SELECT COUNT(DISTINCT login)
FROM student WHERE login LIKE '%@cs'
```

COUNT(DISTINCT login)

10

# AGGREGATES

Output of other columns outside of an aggregate is undefined.

*Get the average GPA of students enrolled in each course.*

```
SELECT AVG(s.gpa), e.cid  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid
```

AVG(s.gpa)	e.cid
3.5	???

The way they fix this

## GROUP BY

Project tuples into subsets and calculate aggregates against each subset.

```
SELECT AVG(s.gpa), e.cid  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid  
GROUP BY e.cid
```

e.sid	s.sid	s.gpa	e.cid
53435	53435	2.25	15-721
53439	53439	2.70	15-721
56023	56023	2.75	15-826
59439	59439	3.90	15-826
53961	53961	3.50	15-826
58345	58345	1.89	15-445

# GROUP BY

Project tuples into subsets and calculate aggregates against each subset.

```
SELECT AVG(s.gpa), e.cid  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid  
GROUP BY e.cid
```

e.sid	s.sid	s.gpa	e.cid
53435	53435	2.25	15-721
53439	53439	2.70	15-721
56023	56023	2.75	15-826
59439	59439	3.90	15-826
53961	53961	3.50	15-826
58345	58345	1.89	15-445

# GROUP BY

Project tuples into subsets and calculate aggregates against each subset.

```
SELECT AVG(s.gpa), e.cid
  FROM enrolled AS e, student AS s
 WHERE e.sid = s.sid
GROUP BY e.cid
```

e.sid	s.sid	s.gpa	e.cid
53435	53435	2.25	15-721
53439	53439	2.70	15-721
56023	56023	2.75	15-826
59439	59439	3.90	15-826
53961	53961	3.50	15-826
58345	58345	1.89	15-445



AVG(s.gpa)	e.cid
2.46	15-721
3.39	15-826
1.89	15-445

# GROUP BY

---

Non-aggregated values in **SELECT** output clause must appear in **GROUP BY** clause.

```
SELECT AVG(s.gpa), e.cid, s.name  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid  
GROUP BY e.cid
```





# GROUP BY

---

Non-aggregated values in **SELECT** output clause must appear in **GROUP BY** clause.

```
SELECT AVG(s.gpa), e.cid, s.name  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid  
GROUP BY e.cid, s.name
```

# HAVING

Filters results based on aggregation computation.

Like a **WHERE** clause for a **GROUP BY**

```
SELECT AVG(s.gpa) AS avg_gpa, e.cid  
FROM enrolled AS e, student AS s  
WHERE e.sid = s.sid  
AND avg_gpa > 3.9  
GROUP BY e.cid
```



This doesn't work because we can't access anything in our aggregations in our where clause because we don't have them yet in our where clause.

The where clause is filtering tuples as we go along and after we do our filtering then we can actually compute aggregation

# HAVING

**Filters** results based on aggregation computation.

Like a **WHERE** clause for a **GROUP BY**

```
SELECT AVG(s.gpa) AS avg_gpa, e.cid  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid  
 GROUP BY e.cid  
 HAVING avg_gpa > 3.9;
```

AVG(s.gpa)	e.cid
3.75	15-415
3.950000	15-721
3.900000	15-826



avg_gpa	e.cid
3.950000	15-721

# STRING OPERATIONS

	String Case	String Quotes
SQL-92	Sensitive	Single Only
Postgres	Sensitive	Single Only
MySQL	Insensitive	Single/Double
SQLite	Sensitive	Single/Double
DB2	Sensitive	Single Only
Oracle	Sensitive	Single Only

**WHERE UPPER(name) = UPPER('KaNyE')** SQL-92

**WHERE name = "KaNyE"** MySQL

# STRING OPERATIONS

**LIKE** is used for string matching.

String-matching operators

- "%" Matches any substring (including empty strings).
- "\_" Match any one character

```
SELECT * FROM enrolled AS e  
WHERE e.cid LIKE '15-%'
```

```
SELECT * FROM student AS s  
WHERE s.login LIKE '%@c_'
```

# STRING OPERATIONS

SQL-92 defines string functions.

→ Many DBMSs also have their own unique functions

Can be used in either output and predicates:

```
SELECT SUBSTRING(name,0,5) AS abbrev_name  
FROM student WHERE sid = 53688
```

```
SELECT * FROM student AS s  
WHERE UPPER(e.name) LIKE 'KAN%'
```

# STRING OPERATIONS

SQL standard says to use **||** operator to concatenate two or more strings together.

```
SELECT name FROM student
WHERE login = LOWER(name) || '@cs'
```

SQL-92

```
SELECT name FROM student
WHERE login = LOWER(name) + '@cs'
```

MSSQL

```
SELECT name FROM student
WHERE login = CONCAT(LOWER(name), '@cs')
```

MySQL

# DATE/TIME OPERATIONS

---

Operations to manipulate and modify **DATE/TIME** attributes.

Can be used in either output and predicates.

Support/syntax varies wildly...

**Demo: Get the # of days since the beginning of the year.**

**There is no standard way to do this even though there is a standard specification**



# OUTPUT REDIRECTION

---

Store query results in another table:

- Table must not already be defined.
- Table will have the same # of columns with the same types as the input.

```
SELECT DISTINCT cid INTO CourseIds SQL-92  
FROM enrolled;
```

```
CREATE TABLE CourseIds (MySQL  
SELECT DISTINCT cid FROM enrolled);
```

# OUTPUT REDIRECTION

---

Insert tuples from query into another table:

- Inner **SELECT** must generate the same columns as the target table.
- DBMSs have different options/syntax on what to do with duplicates.

```
INSERT INTO CourseIds SQL-92  
(SELECT DISTINCT cid FROM enrolled);
```

# OUTPUT CONTROL

## ORDER BY <column\*> [ASC|DESC]

→ Order the output tuples by the values in one or more of their columns.

```
SELECT sid, grade FROM enrolled
WHERE cid = '15-721'
ORDER BY grade
```

sid	grade
53123	A
53334	A
53650	B
53666	D

# OUTPUT CONTROL

## ORDER BY <column\*> [ASC|DESC]

→ Order the output tuples by the values in one or more of their columns.

```
SELECT sid, grade FROM enrolled
WHERE cid = '15-721'
ORDER BY grade
```

sid	grade
53123	A
53334	A
53650	B
53666	D

```
SELECT sid FROM enrolled
WHERE cid = '15-721'
ORDER BY grade DESC, sid ASC
```

sid
53666
53650
53123
53334

# OUTPUT CONTROL

## **LIMIT <count> [offset]**

- Limit the # of tuples returned in output.
- Can set an offset to return a “range”

```
SELECT sid, name FROM student  
WHERE login LIKE '%@cs'  
LIMIT 10
```

```
SELECT sid, name FROM student  
WHERE login LIKE '%@cs'  
LIMIT 20 OFFSET 10
```

**Search results**

# NESTED QUERIES

---

Queries containing other queries.

They are often difficult to optimize.

Inner queries can appear (almost) anywhere in query.

Outer Query → **SELECT name FROM student WHERE  
sid IN (SELECT sid FROM enrolled)** ← Inner Query

# NESTED QUERIES

---

*Get the names of students in '15-445'*

```
SELECT name FROM student  
WHERE ...
```

↑  
“sid in the set of people that take 15-445”

# NESTED QUERIES

---

*Get the names of students in '15-445'*

```
SELECT name FROM student
WHERE ...
      SELECT sid FROM enrolled
      WHERE cid = '15-445'
```



# NESTED QUERIES

---

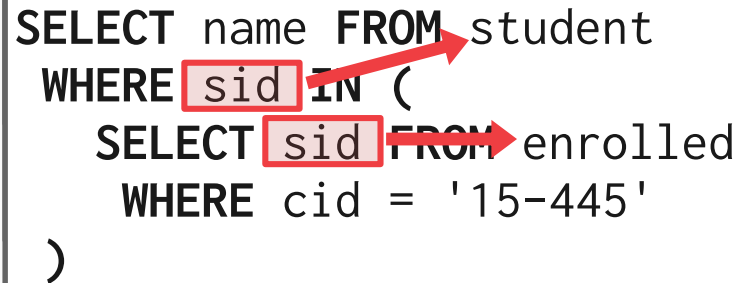
*Get the names of students in '15-445'*

```
SELECT name FROM student
WHERE sid IN (
    SELECT sid FROM enrolled
    WHERE cid = '15-445'
)
```

# NESTED QUERIES

*Get the names of students in '15-445'*

```
SELECT name FROM student
WHERE sid IN (
  SELECT sid FROM enrolled
  WHERE cid = '15-445'
)
```



# NESTED QUERIES

---

**ALL** → Must satisfy expression for all rows in sub-query

**ANY** → Must satisfy expression for at least one row in sub-query.

**IN** → Equivalent to '**=ANY()**'.

**EXISTS** → At least one row is returned.



# NESTED QUERIES

---

*Get the names of students in '15-445'*

```
SELECT name FROM student
WHERE sid = ANY(
  SELECT sid FROM enrolled
  WHERE cid = '15-445'
)
```

# NESTED QUERIES

---

*Get the names of students in '15-445'*

```
SELECT (SELECT S.name FROM student AS S
        WHERE S.sid = E.sid) AS sname
FROM enrolled AS E
WHERE cid = '15-445'
```

## NESTED QUERIES

*Find student record with the highest id that is enrolled in at least one course.*

```
SELECT MAX(e.sid), s.name  
  FROM enrolled AS e, student AS s  
 WHERE e.sid = s.sid;
```



Won't work in SQL-92. This runs in SQLite, but not Postgres or MySQL (v5.7 with strict mode).

# NESTED QUERIES

---

*Find student record with the highest id that is enrolled in at least one course.*

```
SELECT sid, name FROM student  
WHERE ...
```

"Is greater than every other sid"

# NESTED QUERIES

---

*Find student record with the highest id that is enrolled in at least one course.*

```
SELECT sid, name FROM student  
  WHERE sid is greater than every  
    SELECT sid FROM enrolled
```



# NESTED QUERIES

*Find student record with the highest id that is enrolled in at least one course.*

```
SELECT sid, name FROM student
WHERE sid => ALL(
  SELECT sid FROM enrolled
)
```

sid	name
53688	Bieber

# NESTED QUERIES

*Find student record with the highest id that is enrolled in at least one course.*

```
SELECT sid, name FROM student
WHERE sid IN (
  SELECT sid, name FROM student
  WHERE sid IN (
    SELECT MAX(sid) FROM enrolled
  )
)
```

# NESTED QUERIES

*Find student record with the highest id that is enrolled in at least one course.*

```

SELECT sid, name FROM student
WHERE sid IN (
  SELECT sid, name FROM student
  WHERE sid IN (
    SELECT sid FROM enrolled
    ORDER BY sid DESC LIMIT 1
  )
)
  
```

# NESTED QUERIES

*Find all courses that has no students enrolled in it.*

```
SELECT * FROM course
WHERE ...
```

*“with no tuples in the ‘enrolled’ table”*

cid	name
15-445	Database Systems
15-721	Advanced Database Systems
15-826	Data Mining
15-823	Advanced Topics in Databases

sid	cid	grade
53666	15-445	C
53688	15-721	A
53688	15-826	B
53655	15-445	B
53666	15-721	C

# NESTED QUERIES

---

*Find all courses that has no students enrolled in it.*

```
SELECT * FROM course
WHERE NOT EXISTS(
    tuples in the 'enrolled' table
)
```

Can you think of a nested queries as nested for loop?

Ans: Yes but no

# NESTED QUERIES

*Find all courses that has no students enrolled in it.*

```
SELECT * FROM course
WHERE NOT EXISTS(
  SELECT * FROM enrolled
  WHERE course.cid = enrolled.cid
)
```

cid	name
15-823	Advanced Topics in Databases

# NESTED QUERIES

*Find all courses that has no students enrolled in it.*

```
SELECT * FROM course
WHERE NOT EXISTS(
  SELECT * FROM enrolled
  WHERE course.cid = enrolled.cid
)
```

cid	name
15-823	Advanced Topics in Databases

# WINDOW FUNCTIONS

---

Performs a calculation across a set of tuples that related to a single row.

Like an aggregation but tuples are not grouped into a single output tuples.

```
SELECT ... FUNC-NAME(...) OVER (...)
FROM tableName
```



# WINDOW FUNCTIONS

Performs a calculation across a set of tuples that related to a single row.

Like an aggregation but tuples are not grouped into a single output tuples.

```
SELECT ... FUNC-NAME(...) OVER (...)
FROM tableName
```

Aggregation Functions  
Special Functions

# WINDOW FUNCTIONS

Performs a calculation across a set of tuples that related to a single row.

Like an aggregation but tuples are not grouped into a single output tuples.

*How to “slice” up data*  
Can also sort

```
SELECT ... FUNC-NAME(...) OVER (...)
FROM tableName
```

Aggregation Functions  
Special Functions

# WINDOW FUNCTIONS

---

Aggregation functions:

→ Anything that we discussed earlier

Special window functions:

→ **ROW\_NUMBER()** → # of the current row

→ **RANK()** → Order position of the current row.

```
SELECT *, ROW_NUMBER() OVER () AS row_num  
FROM enrolled
```

# WINDOW FUNCTIONS

Aggregation functions:

→ Anything that we discussed earlier

Special window functions:

→ **ROW\_NUMBER()** → # of the current row

→ **RANK()** → Order position of the current row.

sid	cid	grade	row_num
53666	15-445	C	1
53688	15-721	A	2
53688	15-826	B	3
53655	15-445	B	4
53666	15-721	C	5

```
SELECT *, ROW_NUMBER() OVER () AS row_num  
FROM enrolled
```

# WINDOW FUNCTIONS

Aggregation functions:

→ Anything that we discussed earlier

Special window functions:

→ **ROW\_NUMBER()** → # of the current row

→ **RANK()** → Order position of the current row.

sid	cid	grade	row_num
53666	15-445	C	1
53688	15-721	A	2
53688	15-826	B	3
53655	15-445	B	4
53666	15-721	C	5

```
SELECT *, ROW_NUMBER() OVER () AS row_num  
FROM enrolled
```

# WINDOW FUNCTIONS

---

The **OVER** keyword specifies how to group together tuples when computing the window function.

Use **PARTITION BY** to specify group.

```
SELECT cid, sid,  
       ROW_NUMBER() OVER (PARTITION BY cid)  
FROM enrolled  
ORDER BY cid
```

# WINDOW FUNCTIONS

The **OVER** keyword specifies how to group together tuples when computing the window function.

Use **PARTITION BY** to specify group.

cid	sid	row_number
15-445	53666	1
15-445	53655	2
15-721	53688	1
15-721	53666	2
15-826	53688	1

```
SELECT cid, sid,  
       ROW_NUMBER() OVER (PARTITION BY cid)  
FROM enrolled  
ORDER BY cid
```

# WINDOW FUNCTIONS

The **OVER** keyword specifies how to group together tuples when computing the window function.

Use **PARTITION BY** to specify group.

cid	sid	row_number
15-445	53666	1
15-445	53655	2
15-721	53688	1
15-721	53666	2
15-826	53688	1

```
SELECT cid, sid,  
       ROW_NUMBER() OVER (PARTITION BY cid)  
FROM enrolled  
ORDER BY cid
```



# WINDOW FUNCTIONS

---

You can also include an **ORDER BY** in the window grouping to sort entries in each group.

```
SELECT *,  
    ROW_NUMBER() OVER (ORDER BY cid)  
FROM enrolled  
ORDER BY cid
```

# WINDOW FUNCTIONS

*Find the student with the highest grade for each course.*

```
SELECT * FROM (  
  SELECT *,
```

```
    RANK() OVER (PARTITION BY cid  
                  ORDER BY grade ASC)
```

```
  AS rank  
FROM enrolled) AS ranking  
WHERE ranking.rank = 1
```

Group tuples by cid  
Then sort by grade

AS rank

# COMMON TABLE EXPRESSIONS

Provides a way to write auxiliary statements for use in a larger query.

→ Think of it like a temp table just for one query.

Alternative to nested queries and views.

```
WITH cteName AS (  
    SELECT 1  
)  
SELECT * FROM cteName
```

# COMMON TABLE EXPRESSIONS

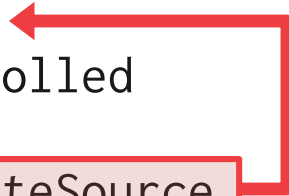
---

You can bind output columns to names before the **AS** keyword.

```
WITH cteName (col1, col2) AS (  
    SELECT 1, 2  
)  
SELECT col1 + col2 FROM cteName
```

# COMMON TABLE EXPRESSIONS

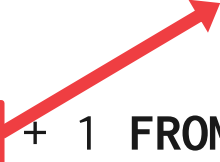
*Find student record with the highest id that is enrolled in at least one course.*

```
WITH cteSource (maxId) AS (  
    SELECT MAX(sid) FROM enrolled  
)  
SELECT name FROM student, cteSource  
WHERE student.sid = cteSource.maxId
```

# CTE – RECURSION

*Print the sequence of numbers from 1 to 10.*

```
WITH RECURSIVE cteSource (counter) AS (  
  (SELECT 1)  
  UNION ALL  
  (SELECT counter + 1 FROM cteSource  
   WHERE counter < 10)  
)  
SELECT * FROM cteSource
```



**Demo: Postgres CTE!**

# CONCLUSION

---

SQL is not a dead language.

You should (almost) always strive to compute your answer as a single SQL statement.



# HOMEWORK #1

---

Write SQL queries to perform basic data analysis on bike-sharing data from SFO.

- Write the queries locally using SQLite.
- Submit them to Gradescope
- You can submit multiple times. We track your best score.

**Due: Monday Sept 10<sup>th</sup> @ 11:59pm**

<https://15445.courses.cs.cmu.edu/fall2018/homework1/>



# NEXT CLASS

---

## Storage Management

